

2021/22

MSc Web and Social Media Analytics

twitter notes

Tutorial Pack 8 (90 minutes)

(To be completed during LW11 tutorial)

LEARNING OBJECTIVES

- To perform exploratory analysis using WordClouds and the Python Counter class
- To practice applying the k-means clustering algorithm to group tweets into themes

LEARNING OUTCOMES

- By the end of this tutorial students will have;
 - Revised key Python programming concepts
 - Practiced interacting with the V2 Twitter API using both the streaming and historic methods.
 - Developed experience writing API queries using rulesets
 - Explored the use of regular expressions to aid formatting of data files
 - Conducted an initial exploratory analysis to identify prominent keywords and phrases
- Applied the k-means approach to group tweets by their TF-IDF scores

RESOURCES AND TOOLS REQUIRED

- Python 3 via repl.it
- Lecture Slides for LW 10
- Rstudio (from apps anywhere)

IMPORTANT:

This pack is designed for you to go at your own speed. At the end of each section there is a series of practice questions and exercises. You should attempt to answer **all questions**.

Any questions you do not complete today should be completed before your next tutorial.

You may find this tutorial pack and the exercises useful preparation for the coursework.

1. Streaming API example (10 mins)

The following code demonstrates how to listen for tweets using the streaming APIv2

twitter #1

```
1 import tweepy
2 import sys
3 import os
4 import codecs
5 import re
6
7 BEARER_TOKEN='[REDACTED]ZD6YjHTUe8H9nkgp6D0UE%3DKsdXQ
8
9 DATA_FILE = "streaming_data.csv"
10 if not os.path.exists(DATA_FILE):
11     os.mknod(DATA_FILE)
12
13 class BasicTwitterListener(tweepy.StreamingClient):
14     def __init__(self, n_tweets):
15         super().__init__(BEARER_TOKEN)
16         self.n_tweets = n_tweets
17
18     def on_tweet(self, tweet):
19         tweet.text = re.sub("[\r\n]+", " ", tweet.text)
20         print(tweet.created_at, tweet.text, tweet.author_id)
21
22         with codecs.open(DATA_FILE, "a") as f:
23             f.write("%s|%s|%s\n" % (tweet.created_at, tweet.text, tweet.author_id))
24
25         self.n_tweets -= 1
26         if self.n_tweets == 0:
27             self.disconnect()
28
29 listener = BasicTwitterListener(10)
30 listener.add_rules(tweepy.StreamRule("news"))
31 listener.filter(expansions=["author_id"], tweet_fields=["created_at"])
```

Study the above code and attempt to answer the following questions.

QNum	Question	Answer
1	On which lines is a file called "streaming_data.csv" created if it doesn't already exist?	
2	Which line numbers contain code that controls the number of tweets that are collected?	
3	What happens when the requested number of tweets are collected?	
4	On which line is data about each tweet written to a file?	
5	When a new tweet arrives, which three fields are saved to the "streaming_data.csv" file?	
6	Which line number controls the tweets that will be listened for?	

2. Historic API example (5 mins)

The following code demonstrates how to listen for tweets using the Historic APIv2

twitter #2

```
main.py x
1 import tweepy
2 import sys
3 import os
4 import codecs
5 import re
6
7 BEARER_TOKEN="AJ5zQfqZD6YjHTUe8H9nkgp6D0UE%3DKsdXQFc
8
9 DATA_FILE = "historic_data.csv"
10 if not os.path.exists(DATA_FILE):
11     os.mknod(DATA_FILE)
12
13 client = tweepy.Client(BEARER_TOKEN)
14
15 response = client.search_recent_tweets(query="news OR weather",
16                                       expansions=["author_id"],
17                                       tweet_fields=['created_at'], max_results=10)
18
19 with codecs.open(DATA_FILE, "a") as f:
20     for tweet in response.data:
21         tweet.text = re.sub("[\r\n]+", " ", tweet.text)
22         print(tweet.created_at, tweet.text, tweet.author_id)
23         f.write("%s| %s| %s\n" % (tweet.created_at, tweet.text, tweet.author_id))
```

Study the above code and attempt to answer the following questions.

QNum	Question	Answer
7	Which two keywords are used to search for relevant tweets?	
8	On which line number can the number of tweets returned be controlled?	
9	How many tweets will appear in the historic_data.csv file if the above code is run twice?	
10	What pattern does the regular expression on line 21 match? Why might it be used here?	

3. Reading back collected tweets (5 mins)

In the two examples seen previously, the tweet text, author id and the date are fields retrieved and saved to a file name specified by the variable DATA_FILE.

The format of this file is as follows: each line represents a single tweet in the form <<created_at>>||<<tweet.text>>||<<tweet.author_id>>. The following function demonstrates how the **tweet text** of all tweets can be extracted from this file.

```

main.py ×
1 import codecs
2
3 def read_tweets_from_file(FILE_NAME):
4     tweets = []
5     with codecs.open(FILE_NAME, "r") as f:
6         for line in f:
7             parts = line.split("||")
8             tweets.append(parts[1])
9     return tweets

```

Study the above code and attempt to answer the following questions.

QNum	Question	Answer
11	How many parameters does the above function accept?	
12	What data type does the above function return?	
13	Explain the purpose of lines 5 and 6	
14	How might the above code be modified to extract only the author_ids?	

Exercise 1: (20 mins)

- Using the streaming API, collect a sample of 100 tweets on a subject of your choice and save them to a file called "historic_data.csv". *Hint:* you will need to modify the sample code provided in the appendix.
- Using the streaming API listen for tweets on a topic of your choice until you have gathered 500 tweets. Save your tweets to "streaming_data.csv"

Exercise 2: (10 mins)

- Create a new repl and upload the data file obtained from exercise 1. Verify you can read in and print your collected tweets using the example code on page 4.
- Modify your code so that tweets are lowercased before being printed i.e., apply a lowercase pre-processing step.

4. Exploratory analysis

An initial exploratory analysis can tell us whether we have applied sufficient preprocessing and indeed chosen relevant keywords. A useful Python class for this purpose is the *Counter* class, which can be found in the *collections* package.

Philip Worrall, PhD, MSc, BAEcon(Hons)

- Tutorial Pack 7 - MSc Web and Social Media Analytics - 4

In the example below I demonstrate how to use the counter class by looping over all the collected tweets, splitting them into their respective words and updating the counter class. At the end of the process I can access the most commonly used words by calling the `most_common()` method.

```
main.py ×
1 import codecs
2 from collections import Counter
3
4 def read_tweets_from_file(FILE_NAME):
5     tweets = []
6     with codecs.open(FILE_NAME, "r") as f:
7         for line in f:
8             parts = line.split("||")
9             tweets.append(parts[1])
10    return tweets
11
12 def cleaned_tweets(FILE_NAME):
13     for line in read_tweets_from_file(FILE_NAME):
14         yield line.lower()
15
16 word_counter = Counter()
17 for line in cleaned_tweets("historic_data.csv"):
18     words = line.split()
19     word_counter.update(words)
20
21 print(word_counter.most_common(10))
```

The output of the counter is shown below. Among the top 10 most used words we can see the words “rt”, “the” and “a”. This indicates a weakness in our preprocessing step since we have not dealt with the issue of stopwords.

```
Console Shell
[('rt', 78), ('the', 59), ('a', 43), ('in', 21), ('to', 27), ('of', 27), ('news', 19), ('and', 18), ('for', 14), ('on', 12)]
```

We have already seen in LW9 that lists of stopwords can be obtained from the NLTK (natural language toolkit) package. In the example below, I have updated my code so that as part of the `cleaned_tweets` function, each word is checked against the list of English stopwords. If a word is found in the list of English stopwords it is excluded.

```

main.py x
1 import codecs
2 from collections import Counter
3 from nltk.corpus import stopwords
4
5 word_counter = Counter()
6 en_stopwords = stopwords.words("english")
7
8 def read_tweets_from_file(FILE_NAME):
9     tweets = []
10    with codecs.open(FILE_NAME, "r") as f:
11        for line in f:
12            parts = line.split("||")
13            tweets.append(parts[1])
14    return tweets
15
16 def cleaned_tweets(FILE_NAME):
17    for line in read_tweets_from_file(FILE_NAME):
18        line = line.lower()
19        words = line.split()
20        words_to_keep = []
21        for word in words:
22            if word not in en_stopwords:
23                words_to_keep.append(word)
24        yield ' '.join(words_to_keep)
25
26 for line in cleaned_tweets("historic_data.csv"):
27     words = line.split()
28     word_counter.update(words)
29
30 print(word_counter.most_common(20))

```

We can see in the output of the top 20 most used words all the stopwords have now been removed. However, this is not to say that all pre-processing has necessarily been completed since we still observe a high incidence of the rt symbol and the ampersand sign relative to other popular terms. Both of which could be candidates for removal in further pre-processing steps.

```

Console  Shell
[('rt', 78), ('news', 19), ('&', 10), ('became', 6), ('|', 6), ('via', 5), ('launch', 5), ('time', 5), ('get', 4), ('today', 4), ('la', 4), ('redmi', 4), ('5g', 4), ('us', 4), ('old', 3), ('see', 3), ('10', 3), ('@flokimooni:', 3)]

```

One way this data can be explored visually is using a word cloud. A word cloud displays all words on a chart in relation to their frequency across a dataset. In Python we can create wordclouds using the wordclouds package. First, we import the wordCloud object from the wordcloud package.


```

main.py ×
1 import codecs
2 from collections import Counter
3 from nltk.corpus import stopwords
4 from wordcloud import WordCloud

```

Next we add a function to create a wordcloud and save it to a file called “wordcloud.png”. This function accepts one parameter, the counter we created earlier that tracks the frequencies of words, to determine the relative sizes of words. In this case, we are drawing the wordcloud based on the top 200 most used words in the dataset.

```

26
27 def make_wordcloud(counts):
28     cloud = WordCloud(width=800, height=400)
29     cloud.generate_from_frequencies(dict(counts.most_common(200)))
30     image = cloud.to_image()
31     image.save("wordcloud.png")
32

```

To draw the wordcloud we must call the make_wordcloud function, passing the word_counter variable as a parameter, at the end of the program when all words have been counted.

```

37 make_wordcloud(word_counter)

```

Table A. Python code used to create a wordcloud

```

import codecs from collections
import Counter from nltk.corpus
import stopwords
from wordcloud import WordCloud

word_counter = Counter()
en_stopwords = stopwords.words("english")

def read_tweets_from_file(FILE_NAME):
    tweets = []
    with
    codecs.open(FILE_NAME, "r") as f:
        for
        line in f:
            parts = line.split("|")
            tweets.append(parts[1])
    return tweets

def cleaned_tweets(FILE_NAME):
    for line in read_tweets_from_file(FILE_NAME):
        line = line.lower()
        words =
        line.split()
        words_to_keep = []
        for word in words:
            if word not
            in en_stopwords:
                words_to_keep.append(word)
        yield ' '.join(words_to_keep)

def make_wordcloud(counts):
    cloud =
    WordCloud(width=800, height=400)
    cloud.generate_from_frequencies(dict(counts.most_common(200)))
    image = cloud.to_image()
    image.save("wordcloud.png")

for line in cleaned_tweets("historic_data.csv"):
    words = line.split()
    word_counter.update(words)

make_wordcloud(word_counter)

```



```
> library(tm)
> library(cluster)
```

As a case study I am going to use the previously collected tweets surrounding the release of the Gravity movie. They are contained in a text file called “gravity-filmtweets.txt” that is placed in the working directory of my R installation.

In the following code I read in the tweets into a new document corpus and perform some initial pre-processing. If the data has already been processed when it was collected (i.e. from Python) these additional steps are not strictly necessary.

Table 2 Example code use to load in the gravity film tweets into R

```
text <- readLines("gravity-film-tweets.txt", encoding="UTF-8")
doc.vec <- VectorSource(text) doc.corpus <-
Corpus(doc.vec)
doc.corpus <- tm_map(doc.corpus, tolower)
doc.corpus <- tm_map(doc.corpus, removeWords,
stopwords('english')) doc.corpus <- tm_map(doc.corpus,
removePunctuation) doc.corpus <- tm_map(doc.corpus,
removeNumbers) doc.corpus <- tm_map(doc.corpus, stripWhitespace)
```

In the next step I convert the corpus (collection of text documents) into a document term matrix. This is essentially the bag of words model we saw in LW9.

```
dtm <- DocumentTermMatrix(doc.corpus)
dtm
```

From the output below we see we have a total of 3606 tweets that use 5008 unique words (our dimensions). The longest word used is 31 characters. The high level of sparsity implies that individual tweets use only a tiny proportion of available words (as expected)

```
<<DocumentTermMatrix (documents: 3606, terms: 5008)>>
Non-/sparse entries: 29706/18029142
Sparsity             : 100%
Maximal term length: 31
weighting            : term frequency (tf)
> |
```

Since high levels of sparsity can weaken statistical relationships, we next remove terms that don't appear in at least 1% of our dataset.

```
dtm <- removeSparseTerms(dtm, 0.99) dtm
```

The output below shows that following a removal of sparse terms, the number of unique dimensions has dropped dramatically to just 87. Thus 99% of all tweets contain one or more of these dimensions.

```
> dtm <- removesparseTerms(dtm, 0.99)
> dtm
<<DocumentTermMatrix (documents: 3606, terms: 87)>>
Non-/sparse entries: 15704/298018
sparsity           : 95%
Maximal term length: 16
weighting          : term frequency (tf)
```

We can access these remaining 87 words using the `dimnames` property of the document term matrix.

```
dtm$dimnames
```

And obtain the following list of terms

```
$Terms
[1] "film"           "gravity"         "just"           "space"          "good"           "better"
[7] "finally"        "incredible"     "watched"        "like"           "review"         "ever"
[13] "seen"           "stressful"      "much"           "really"         "watching"       "acclaimed"
[19] "critically"     "deleted"        "photo"          "scene"          "boring"         "pretty"
[25] "well"           "brilliant"      "shit"           "wow"            "one"            "think"
[31] "bullock"        "got"            "sandra"         "see"            "amazing"        "time"
[37] "load"           "made"           "tonight"        "intense"        "watch"          "great"
[43] "now"            "story"          "oscars"         "won"            "best"           "amp"
[49] "last"           "night"          "worst"          "movie"          "awesome"        "awards"
[55] "many"           "can"            "effects"        "real"           "saw"            "know"
[61] "award"          "elliottexplicit" "..."          "ittybittymanatee" "get"            "cuaron"
[67] "alfonso"        "moving"         "astronauts"     "winning"        "years"          "oscar"
[73] "still"          "work"           "wins"           "gravity"         "oscarwinning"   "beyond"
[79] "variety"        "cameraman"      "chief"          "collects"        "honour"         "wor"
[85] "exclusively"    "talked"         "georgelineker"
```

Clearly the film's name is an important term, but what other terms are most often used in the same tweet. For this analysis we can use a function called `findAssocs()`, which identifies all terms that have a higher than specified correlation with the given term.

```
> findAssocs(dtm, "gravity", 0.5)
$gravity
  cameraman    collects    honour    chief oscarwinning    wor    award
      0.60         0.60         0.60      0.59         0.58         0.53         0.51
```

Previously, when we removed sparse terms, we deleted all terms that didn't appear in at least 1% of our collected tweets. As a result, we may now have tweets in our collection that contain a row full of zeros (since all their words were deleted). We will now need to remove these documents since they don't use any of the remaining terms. A simple sum of each row to check that it adds to a positive number will suffice.

```
rowTotals <- apply(dtm, 1, sum) dtm <-
dtm[rowTotals > 0, ]
```

```
dtm
```

From the output below we found approximately $(3606-3463 = 143)$ of such documents.


```
> rowTotals <- apply(dtm, 1, sum)
> dtm <- dtm[rowTotals > 0, ]
> dtm
<<DocumentTermMatrix (documents: 3463, terms: 87)>>
Non-/sparse entries: 15704/285577
Sparsity           : 95%
Maximal term length: 16
weighting          : term frequency (tf)
> |
```

Next we create the TF-IDF matrix to replace the term frequencies and calculate the normalised Euclidean distance matrix. The distance matrix measures how far each document is from all other documents as the “crow flies”. This measure is used by the k-means algorithm to determine whether any two documents should be placed in the same cluster.

```
dtm_tfidf <- weightTfIdf(dtm) m
<- as.matrix(dtm_tfidf)
rownames(m) <- 1:nrow(m)
norm_eucl <- function(m) m/apply(m, 1, function(x) sum(x^2)^.5) m_norm
<- norm_eucl(m)
```

```
> dtm_tfidf
<<DocumentTermMatrix (documents: 3463, terms: 87)>>
Non-/sparse entries: 15704/285577
Sparsity           : 95%
Maximal term length: 16
weighting          : term frequency - inverse document frequency (normalized) (tf-idf)
> |
```

The Euclidean distance comes from Pythagoras’s famous theorem. In the twodimensional case, the Euclidean distance between points A and B is given by the vector k. It is found using the square root of the sum of the squares of both dimensions i and j (i.e. $A^2+B^2 = C^2$)



The k-means method can now be applied. As a starting point we choose the number of clusters to be the square root of the number of text documents in our dataset. This can be fine tuned later using an elbow approximation or other suitable method.

```
results <- kmeans(m_norm, as.integer(floor(length(text)^0.5)))
```

The results variable can now be inspected to check the results following model fitting. In this case we observe a ratio of the BSS to the TSS of 71.8% which is a reasonable initial fit.

```
within cluster sum of squares by cluster:
[1]  5.064932 19.791356 18.779285  8.507911 16.153930 15.288170 20.386066  6.930494  3.008845 13.954560  9.863083
[12] 15.579038 12.652347 14.283637 14.053072  1.218796  5.180480  1.542760 31.256611 12.849378 10.451014 12.253511
[23]  1.715865 11.321460 14.371840 24.080205  5.337156  9.975769 26.660443 11.395404 21.973306 14.643056 11.745965
[34] 14.279035 33.058659 15.872173  9.366811  7.833001  8.481059 12.534676  8.881359  8.925181 17.115030 12.252177
[45]  4.434073 15.732449 43.127863 20.252757 61.775801 11.171846 15.878483 12.593047 28.976306 17.103279 33.939210
[56] 24.201758  3.861506 29.560081 11.525134 11.058511
(between_ss / total_ss = 71.7 %)
```

By inspection of the size property, we can see that a large proportion of tweets were placed in clusters 49 (600 tweets), cluster 12 (245 tweets) and cluster 29 (181 tweets). Cluster 16 has the lowest number of tweets assigned (5 tweets)

```
> results$size
[1] 11 38 40 34 31 50 87 155 20 30 37 245 70 32 39 5 20 8 67 41 38 34 38 37 32 61 10 30 181
[30] 42 43 44 32 29 74 39 29 25 37 55 17 23 46 27 11 36 84 53 600 33 39 54 120 41 42 76 72 43
[59] 39 37
>
```

Inspecting the centres of the largest clusters we observe a potential theme of discussion around “critically”, “acclaimed”, “deleted”, “photo” and “scene” in cluster 29.


```

> results$centers[29,]
      film      gravity      just      space      good      better      finally
0.011511707 0.010817606 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
incredible      watched      like      review      ever      seen      stressful
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
much      really      watching      acclaimed      critically      deleted      photo
0.000000000 0.000000000 0.000000000 0.409072102 0.409072102 0.409072102 0.335885633
scene      boring      pretty      well      brilliant      shit      wow
0.395236701 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
one      think      bullock      got      sandra      see      amazing
0.000000000 0.000000000 0.002268328 0.000000000 0.004565600 0.002317355 0.000000000
time      load      made      tonight      intense      watch      great
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
now      story      oscars      won      best      amp      last
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
night      worst      movie      awesome      awards      many      can
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
effects      real      saw      know      award      elliotexplicit      -
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.263822150 0.017764836
ittybittymantee      get      cuaron      alfonso      moving      astronauts      winning
0.107027838 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
years      oscar      still      work      wins      'gravity'      oscarwinning
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
beyond      variety      cameraman      chief      collects      honour      wor
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
exclusively      talked      georgelineker      0.000000000
0.000000000 0.000000000 0.000000000
> |

```

Cluster 12 has high relative usage of the terms “cameraman”, “oscarwinning”, “honour”, “won”, “wins” and “award”.

```

> results$centers[12,]
      film      gravity      just      space      good      better      finally
0.009608966 0.017697969 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
incredible      watched      like      review      ever      seen      stressful
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
much      really      watching      acclaimed      critically      deleted      photo
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
scene      boring      pretty      well      brilliant      shit      wow
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
one      think      bullock      got      sandra      see      amazing
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
time      load      made      tonight      intense      watch      great
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
now      story      oscars      won      best      amp      last
0.000000000 0.001932043 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
night      worst      movie      awesome      awards      many      can
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
effects      real      saw      know      award      elliotexplicit      -
0.000000000 0.000000000 0.000000000 0.000000000 0.256982045 0.000000000 0.000000000
ittybittymantee      get      cuaron      alfonso      moving      astronauts      winning
0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000 0.000000000
years      oscar      still      work      wins      'gravity'      oscarwinning
0.000000000 0.000000000 0.000000000 0.046304929 0.272860070 0.000000000 0.294634023
beyond      variety      cameraman      chief      collects      honour      wor
0.000000000 0.000000000 0.604691568 0.303607703 0.305278109 0.305278109 0.249675957
exclusively      talked      georgelineker      0.000000000
0.000000000 0.000000000 0.000000000
> |

```

Cluster 49 has very high average usage of the words “film” and “gravity” but almost no other terms. This might reflect tweets just announcing the film rather than any discussion of it.

```

> results$centers[49,]
      film      gravity      just      space      good      better      finally
0.6716713 0.6677559 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
incredible      watched      like      review      ever      seen      stressful
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
much      really      watching      acclaimed      critically      deleted      photo
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
scene      boring      pretty      well      brilliant      shit      wow
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
one      think      bullock      got      sandra      see      amazing
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
time      load      made      tonight      intense      watch      great
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
now      story      oscars      won      best      amp      last
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
night      worst      movie      awesome      awards      many      can
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
effects      real      saw      know      award      elliotexplicit      -
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
ittybittymantee      get      cuaron      alfonso      moving      astronauts      winning
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
years      oscar      still      work      wins      'gravity'      oscarwinning
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
beyond      variety      cameraman      chief      collects      honour      wor
0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000 0.0000000
exclusively      talked      georgelineker      0.0000000
0.0000000 0.0000000 0.0000000
>

```

Looking at cluster 16, which had the lowest number of tweets assigned, we see low levels of average term usage except for the terms “really” and “watching”.

```

> results$centers[16,]
      film      gravity      just      space      good      better      finally
0.01821448 0.01381692 0.08234453 0.00000000 0.07965265 0.00000000 0.00000000
incredible watched      like      review      ever      seen      stressful
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
much      really      watching      acclaimed      critically      deleted      photo
0.00000000 0.64816058 0.52947472 0.00000000 0.00000000 0.00000000 0.00000000
scene      boring      pretty      well      brilliant      shit      wow
0.00000000 0.00000000 0.00000000 0.00000000 0.12928894 0.00000000 0.00000000
one      think      bullock      got      sandra      see      amazing
0.10035138 0.00000000 0.08761103 0.00000000 0.08761103 0.00000000 0.00000000
time      load      made      tonight      intense      watch      great
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
now      story      oscar      won      best      amp      last
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
night      worst      movie      awesome      awards      many      can
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
effects      real      saw      know      award      elliotexplicit      ...
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
ittybittymanatee      get      cuaron      alfonso      moving      astronauts      winning
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
years      oscar      still      work      wins      'gravity'      oscarwinning
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
beyond      variety      cameraman      chief      collects      honour      wor
0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000 0.00000000
exclusively      talked      georgelineker
0.00000000 0.00000000 0.00000000
> |

```

Table 3 Sample code for k-means in R

```

library(tm)
library(cluster)

text <- readLines("gravity-film-tweets.txt", encoding="UTF-8")
doc.vec <- VectorSource(text) doc.corpus <-
Corpus(doc.vec)
doc.corpus <- tm_map(doc.corpus, tolower)
doc.corpus <- tm_map(doc.corpus, removeWords,
stopwords('english')) doc.corpus <- tm_map(doc.corpus,
removePunctuation) doc.corpus <- tm_map(doc.corpus,
removeNumbers) doc.corpus <- tm_map(doc.corpus, stripWhitespace)

dtm <- DocumentTermMatrix(doc.corpus) dtm

dtm <- removeSparseTerms(dtm, 0.99) dtm

rowTotals <- apply(dtm, 1, sum) dtm
<- dtm[rowTotals > 0, ]
dtm

dtm_tfxfidf <- weightTfIdf(dtm) m
<- as.matrix(dtm_tfxfidf)
rownames(m) <- 1:nrow(m)
norm_eucl <- function(m) m/apply(m, 1, function(x) sum(x^2)^.5) m_norm
<- norm_eucl(m)

results <- kmeans(m_norm, as.integer(floor(length(text)^0.5)))

```


Exercise 4: (15 mins)

- Using the sample code as a guide, apply the k-means approach to the dataset collected using the streaming API as part of exercise 1.
 - Identify the initial number of unique dimensions (words) that are part of your dataset
 - After removing 1% sparse terms, how many word dimensions remain?
 - Take a screenshot of the remaining terms extracted from the document term matrix.
 - Comment on the ratio of the BSS/TSS ratio. Experiment with different numbers of clusters (+/- 2), how does the BSS/TSS ratio change?
 - Inspect the clusters with the highest number of tweets assigned and examine their word usage.

Appendix:

Table 4 Streaming Code

```
import tweepy
import sys
import os
import codecs
import re

BEARER_TOKEN="INSERT BEARER TOKEN HERE"

DATA_FILE = "streaming_data.csv"
if not os.path.exists(DATA_FILE):
    os.mknod(DATA_FILE)

class BasicTwitterListener(tweepy.StreamingClient):
    def __init__(self, n_tweets):
        super().__init__(BEARER_TOKEN)
        self.n_tweets = n_tweets

    def on_tweet(self, tweet):
        tweet.text = re.sub("[\r\n]+", " ", tweet.text)
        print(tweet.created_at, tweet.text, tweet.author_id)

        with codecs.open(DATA_FILE, "a") as f:
            f.write("%s||%s||%s\n" % (tweet.created_at, tweet.text, tweet.author_id))

        self.n_tweets -= 1
        if self.n_tweets == 0:
            self.disconnect()

listener = BasicTwitterListener(10)
listener.add_rules(tweepy.StreamRule("news"))
listener.filter(expansions=["author_id"], tweet_fields=["created_at"])
```

Table 5 Historic Code

```
import tweepy
import sys
import os
import codecs
import re

BEARER_TOKEN="INSERT BEARER TOKEN HERE"

DATA_FILE = "historic_data.csv"
if not os.path.exists(DATA_FILE):
    os.mknod(DATA_FILE)

client = tweepy.Client(BEARER_TOKEN)

response = client.search_recent_tweets(query="news OR weather",
                                       expansions=["author_id"],
                                       tweet_fields=['created_at'], max_results=100)

with codecs.open(DATA_FILE, "a") as f:
    for tweet in response.data:
        tweet.text = re.sub("[\r\n]+", " ", tweet.text)
        print(tweet.created_at, tweet.text, tweet.author_id)
        f.write("%s||%s||%s\n" % (tweet.created_at, tweet.text, tweet.author_id))
```

END